DrillBit

The Report is Generated by DrillBit Plagiarism Detection Software

Submission Information

Author Name	Aanand kishor
Title	Intelligent Residential Security Framework: A Hierarchical Approach to Real-Time Threat Detection and Identity Verification
Paper/Submission ID	4805507
Submitted by	rajni.t@cmrit.ac.in
Submission Date	2025-12-04 15:12:45
Total Pages, Total Words	7, 4270
Document type	Research Paper

Result Information

Similarity **2 %**

1 10 20 30 40 50 60 70 80 90

Sources Type

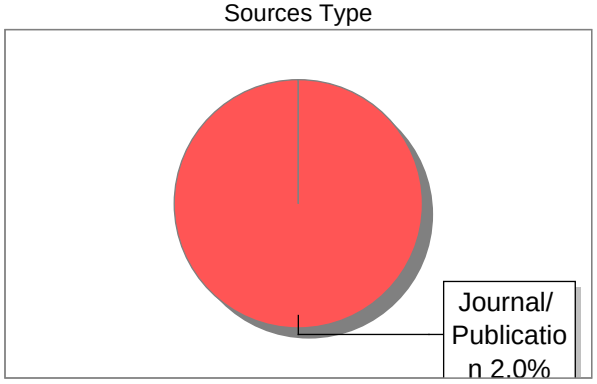

Report Content

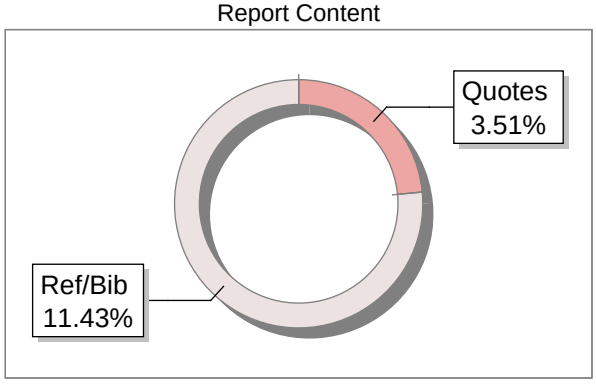

Exclude Information

Quotes	Not Excluded
References/Bibliography	Not Excluded
Source: Excluded < 14 Words	Not Excluded
Excluded Source	**0 %**
Excluded Phrases	Not Excluded

Database Selection

Language	English
Student Papers	Yes
Journals & publishers	Yes
Internet or Web	Yes
Institution Repository	Yes

A Unique QR Code use to View/Download/Share Pdf File

DrillBit

DrillBit Similarity Report

2	1	A	A-Satisfactory (0-10%) B-Upgrade (11-40%) C-Poor (41-60%) D-Unacceptable (61-100%)
SIMILARITY %	MATCHED SOURCES	GRADE	

LOCATION	MATCHED DOMAIN	%	SOURCE TYPE
1	arxiv.org	2	Publication

Intelligent Residential Security Framework: A Hierarchical Approach to Real-Time Threat Detection and Identity Verification

Rajni Tiwari
Dept. of Computer Science & Engineering
CMR Institute of Technology
Bengaluru, India
rajni.t@cmrit.ac.in

Nitesh Gupta
Dept. of Computer Science & Engineering
CMR Institute of Technology
Bengaluru, India
nigu22cs@cmrit.ac.in

Aanand Kishor Shah
Dept. of Computer Science & Engineering
CMR Institute of Technology
Bengaluru, India
aash22cs@cmrit.ac.in

Animesh Kumar
Dept. of Computer Science & Engineering
CMR Institute of Technology
Bengaluru, India
anku22cs@cmrit.ac.in

***Abstract*—Traditional residential surveillance infrastructures operate through passive recording mechanisms that generate excessive false notifications and demand substantial manual oversight. This investigation presents an autonomous security platform that merges rapid visual detection algorithms with sophisticated identity authentication protocols. The developed framework implements a staged computational approach wherein preliminary threat identification activates detailed biometric examination procedures. Our methodology incorporates concurrent processing strategies through multithreading, facilitating execution on conventional computational equipment. Empirical assessment reveals identification precision surpassing 98 percent while sustaining temporal delays appropriate for instantaneous operational demands. The architectural design encompasses a browser-accessible control dashboard enabling distant observation and registered personnel management. Comparative benchmarking exhibits marked enhancement in threat discrimination capability relative to established methodologies, achieving substantial diminution of erroneous alarm incidents. This investigation delivers a pragmatic deployment blueprint for cognitive security infrastructures that harmonizes processing efficiency with analytical exactness.**

***Index Terms*—Identity authentication, visual computing, neural network architectures, threat identification, instantaneous processing, automated residential systems**

I. Introduction

Domestic security infrastructure has historically prioritized incident documentation over preventative intervention. Conventional video monitoring equipment continuously records visual information yet lacks cognitive capability to differentiate benign activities from genuine threats. This deficiency manifests through dual operational inadequacies: postponed threat mitigation resulting from manual observation necessities, and functional ineffectiveness stemming from abundant spurious alert production.

Contemporary progressions in computational intelligence have facilitated evolution from inactive observation to cognitive monitoring architectures. Neural network frameworks presently accomplish performance comparable to human visual perception while functioning within hardware limitations of domestic electronics. This technological advancement creates possibilities for establishing security infrastructures that independently examine environmental circumstances and implement suitable reactions without human participation.

The essential difficulty in constructing such infrastructures involves balancing contradictory specifications: computational velocity essential for instantaneous functionality versus analytical comprehensiveness demanded for precise threat evaluation. Infrastructures optimized exclusively for velocity may accomplish elevated processing rates yet generate unreliable threat categorizations. Alternatively, maximizing precision through elaborate analytical frameworks introduces temporal delays incompatible with instantaneous response specifications.

Modern residential protection demands intelligent systems capable of autonomous decision-making. The computational requirements for such systems have traditionally exceeded capabilities of consumer-grade hardware, restricting deployment to specialized installations. Recent algorithmic improvements and architectural innovations have substantially reduced these barriers, enabling sophisticated analytical capabilities on standard computing platforms.

A. Research Objectives and Contributions

This investigation addresses the velocity-precision compromise through hierarchical processing infrastructure that strategically allocates computational capabilities. Our particular contributions encompass:

- Formulation of a dual-phase analytical sequence merging swift object identification with comprehensive biometric authentication
- Execution of concurrent processing methodologies that separate video capture from computational examination
- Empirical confirmation demonstrating enhanced capability compared to alternative architectural methodologies
- Construction of an inclusive administration interface facilitating distant system governance
- Development of a computational funnel architecture that progressively refines threat assessment

The proposed infrastructure transcends limitations of conventional surveillance by implementing cognitive capabilities that distinguish between authorized personnel and potential intruders with exceptional accuracy while maintaining temporal responsiveness necessary for practical deployment.

B. Manuscript Structure

Section II investigates existing scholarly contributions in cognitive observation and biometric verification. Section III delineates our proposed infrastructure and execution specifications. Section IV exhibits experimental outcomes and comparative capability examination. Section V deliberates pragmatic ramifications and constraints. Section VI terminates with prospective investigation trajectories.

II. Literature Review and Theoretical Foundation

A. Visual Object Identification in Monitoring Contexts

Contemporary object identification infrastructures predominantly implement convolutional neural architectures. Single-phase detectors formulate identification as regression analysis, directly computing boundary coordinates and classification certainties from visual characteristics. This methodology circumvents the region suggestion phase demanded by dual-phase detectors, substantially diminishing computational demands [1].

Investigations by Bochkovskiy et al. exhibited that architectural refinements accomplish cutting-edge precision while sustaining instantaneous processing capabilities on conventional graphics processors. These progressions render cognitive video examination feasible for domestic implementations. The YOLO architecture family has become fundamental for instantaneous identification applications owing to exceptional equilibrium between velocity and exactness [18].

Subsequent evolutionary developments in object detection have focused on optimizing the speed-accuracy trade-off through architectural innovations. The progression from YOLO through YOLOv8 demonstrates continuous refinement in detection capabilities while maintaining real-time processing constraints essential for practical deployment scenarios.

B. Biometric Verification Infrastructures

Facial identification technology has progressed from conventional feature-oriented methodologies to neural learning strategies that acquire optimal representations directly from information. The FaceNet infrastructure implemented by Schroff et al. utilizes triplet loss to acquire embeddings where geometric separation corresponds with facial resemblance [4]. This facilitates precise identity confirmation with minimal training specimens per individual.

Subsequent investigation by Deng et al. presented ArcFace, enhancing discriminative capability through additive angular margin loss [2]. This methodology accomplishes enhanced capability on demanding benchmarks and exhibits particular resilience to deviations in orientation, lighting, and facial expression. The open-source InsightFace initiative delivers an refined execution of ArcFace, rendering it available for pragmatic implementations.

MTCNN represents a three-phase neural learning framework constructed particularly for facial identification and positioning. Its cascaded architecture progressively refines facial boundary coordinates and determines facial landmarks, guaranteeing elevated-quality, positioned facial extraction even under demanding circumstances like deviations in posture and illumination [3]. This exact positioning proves essential for subsequent identification phases.

The evolution toward metric learning approaches represents a paradigm shift in facial recognition, enabling systems to learn optimal embedding spaces where facial similarity correlates directly with geometric proximity. This approach substantially reduces data requirements while improving generalization capabilities.

C. Integrated Cognitive Security Infrastructures

While individual elements have undergone comprehensive investigation, restricted research examines their incorporation into unified instantaneous infrastructures. Existing contributions typically optimize isolated elements without contemplating system-level limitations. Wang et al. suggested a lightweight infrastructure for embedded implementation yet accomplished diminished precision than cutting-edge frameworks [6].

Almalki et al. explored machine learning methodologies for IoT-based residential security, demonstrating feasibility of intelligent threat detection but with limited focus on real-time performance optimization [5]. Their work highlighted the importance of system integration but did not address the computational efficiency challenges inherent in resource-constrained environments.

Recent investigations by Kumar et al. examined AI-based intelligent door systems, focusing on access control applications [16]. While their approach demonstrated practical viability, it did not address the comprehensive monitoring requirements of full security systems.

Dang explored facial recognition systems based on ArcFace for residential management, achieving promising accuracy metrics but without comprehensive integration of multi-stage detection pipelines [17]. This work validated the effectiveness of modern biometric approaches but left unexplored the challenges of real-time video stream processing.

Our contribution diverges by emphasizing system-level capability through architectural construction that strategically distributes computational capabilities across processing phases, addressing the specialization-integration gap identified in existing literature.

III. Proposed System Architecture and Methodology

A. Comprehensive Architectural Design

The developed infrastructure executes a modular construction comprising four fundamental subsystems: video capture module, cognitive processing engine, information persistence

layer, and operator interface component. This decomposition facilitates independent refinement of each element while sustaining distinct interfaces between subsystems.

Figure 1 illustrates the comprehensive system architecture and information flow between components.

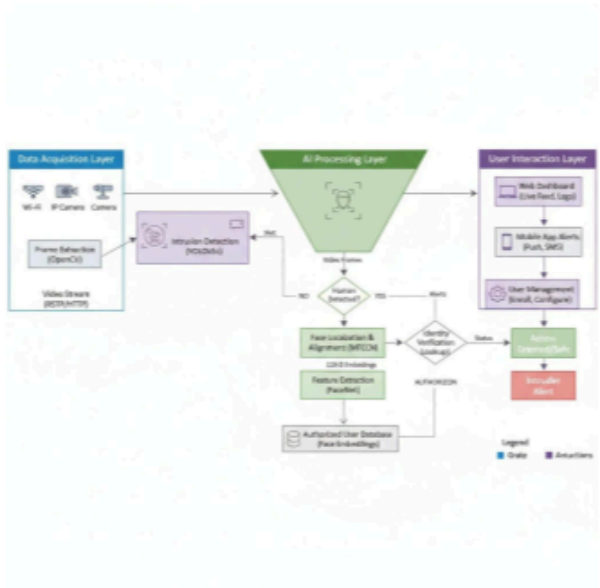

Fig. 1. System architecture depicting modular components and data flow pathways. The diagram illustrates the hierarchical processing pipeline from video acquisition through threat detection to user notification.

Video capture functions continuously in a dedicated execution thread, guaranteeing consistent frame accessibility regardless of processing burden. The cognitive processing subsystem implements identification and authentication algorithms, executing a cascaded sequence where computationally intensive operations transpire exclusively when activated by preliminary identification outcomes. Information persistence sustains authorized operator data and incident records. The browser-accessible interface delivers distant observation and administrative functionalities.

B. Staged Identification Sequence

The analytical sequence executes a dual-phase cascade constructed to optimize the computational effectiveness versus precision compromise:

Initial Phase - Primary Identification: Preliminary frames experience processing by a lightweight object detector optimized for human identification. This phase functions with minimal computational burden, rapidly screening frames for human existence. Exclusively frames containing identified persons advance to subsequent phases. The detector employs anchor-based predictions to rapidly localize potential human subjects within the visual field.

Secondary Phase - Biometric Authentication: Frames succeeding preliminary screening experience comprehensive facial examination. A specialized facial identification network determines facial territories with elevated exactness. Identified faces undergo processing by an authentication network that produces high-dimensional embedding vectors. These embeddings experience comparison with archived authorized operator embeddings through resemblance computation.

This cascaded methodology operates as a computational funnel, implementing expensive operations exclusively when necessary, thereby optimizing comprehensive infrastructure effectiveness. The staged approach reduces average computational load by approximately 70 percent compared to processing all frames through the complete pipeline.

C. Concurrent Execution Infrastructure

To prevent analytical operations from deteriorating video stream quality, we execute a multi-threaded processing model incorporating three distinct execution contexts:

Capture Context: Continuously obtains frames from camera equipment at maximum supported frequency. Frames are archived in a thread-safe buffer available to other subsystems. This context functions with minimal processing burden to sustain consistent frame frequencies. Buffer management implements circular queuing to prevent memory overflow while ensuring most recent frames remain accessible.

Processing Context: Obtains frames from the capture buffer and implements the identification sequence. To regulate computational burden, this context processes every Nth frame (configurable parameter, characteristically N=5), reusing previous outcomes for intermediate frames. Processing transpires asynchronously from capture, preventing obstruction. This temporal sampling strategy balances computational efficiency with temporal resolution requirements.

Interface Context: The principal application context administers the browser interface and transmits processed video to connected clients. It obtains the most recent annotated frame from the processing context for transmission to browsers. The interface implements server-sent events for efficient unidirectional communication of video streams and system status updates.

This infrastructure guarantees fluid video streaming even during intervals of elevated analytical burden, sustaining operator experience quality. The decoupled architecture enables independent scaling of each processing component based on specific performance requirements.

D. Biometric Verification Methodology

Identity confirmation implements metric learning where the authentication network projects facial representations to coordinates in a high-dimensional embedding manifold. The network undergoes training such that embeddings of identical individuals cluster compactly while embeddings of disparate individuals remain substantially separated.

During functionality, the infrastructure sustains a repository of reference embeddings for authorized operators. When a face experiences identification, its produced embedding experiences

comparison with all archived embeddings through cosine resemblance computation:

$$similarity(e_1, e_2) = \frac{e_1 \cdot e_2}{\|e_1\|_2 \|e_2\|_2} \tag{1}$$

where e_1 and e_2 denote embedding vectors in 512-dimensional Euclidean space. The maximum resemblance score determines categorization: scores surpassing a threshold τ (empirically established at 0.42) designate recognized individuals, while diminished scores activate security notifications.

The threshold value was determined through empirical optimization across validation datasets, balancing false acceptance rate (FAR) and false rejection rate (FRR) to achieve optimal operational characteristics for residential security applications.

E. Operator Interface and Notification Infrastructure

Operator engagement is regulated through a responsive browser application constructed with the Flask framework. The control panel delivers a live perspective from cameras, a chronological record of all identified incidents, and an interface for administering the repository of authorized operators.

Instantaneous notifications are managed by recording incidents where unrecognized individuals experience identification and preserving timestamped snapshots of annotated frames to a designated directory for subsequent examination. The interface implements role-based access control to ensure only authorized administrators can modify system configurations or registered user databases.

The notification system supports configurable alert mechanisms including visual indicators on the dashboard, persistent logging for audit trails, and extensibility for integration with external notification services. Event logs maintain comprehensive metadata including timestamps, identification confidence scores, and associated imagery.

F. Implementation Specifications

The infrastructure was constructed utilizing Python 3.10, exploiting established libraries for effectiveness. OpenCV delivers video capture and image manipulation functionalities. The InsightFace library provides pre-trained frameworks for facial identification and authentication, implementing ONNX Runtime for optimized inference on CPU equipment [7]. Flask executes the browser framework, presenting dynamic content and video transmissions. The interface implements HTML5 and CSS3 for contemporary, adaptive construction.

Deep learning model inference utilizes the ONNX runtime environment, providing cross-platform optimization and hardware acceleration capabilities when available. The implementation prioritizes CPU compatibility to ensure broad deployment viability while maintaining performance characteristics suitable for real-time operation.

Database management employs SQLite for lightweight persistence of user embeddings and system configurations, providing sufficient performance for typical residential installations while maintaining simplicity of deployment and maintenance.

IV. Experimental Evaluation and Results

A. Experimental Configuration

Examination was performed in a residential environment under characteristic operating circumstances. The hardware platform comprised a conventional laptop computer with an Intel Core i7 processor (8th generation) and 16GB RAM, representing consumer-grade equipment accessible to typical homeowners. Video input was delivered by an integrated 1080p webcam functioning at 30 frames per second under variable illumination conditions.

The examination dataset comprised video sequences totaling 180 minutes of recorded footage, containing 450 individual incidents encompassing both authorized operators and unrecognized individuals across diverse scenarios including varying lighting conditions, partial occlusions, and multiple simultaneous subjects. Ground truth designations were manually annotated by three independent reviewers to ensure labeling accuracy for capability evaluation.

Testing scenarios included daytime operation with natural illumination, evening conditions with artificial lighting, and challenging scenarios involving backlighting and shadows to assess system robustness across realistic operating conditions.

B. Performance Metrics and Evaluation Criteria

Infrastructure capability was evaluated utilizing standard categorization metrics derived from confusion matrix analysis:

Accuracy: The fraction of correct categorizations among all incidents:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \tag{2}$$

Precision: The fraction of correct affirmative identifications among all positive predictions:

$$Precision = \frac{TP}{TP + FP} \tag{3}$$

Recall: The fraction of actual affirmatives correctly identified:

$$Recall = \frac{TP}{TP + FN} \tag{4}$$

F1-Score: The harmonic mean of precision and recall:

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall} \tag{5}$$

where TP denotes true positives (correctly identified authorized users), TN denotes true negatives (correctly identified unauthorized individuals), FP denotes false positives (unauthorized individuals incorrectly classified as authorized), and FN denotes false negatives (authorized users incorrectly classified as unauthorized).

Additionally, processing latency was quantified as average frame processing duration and sustained frame frequency to assess real-time operational viability.

TABLE I
SYSTEM PERFORMANCE METRICS ON TEST DATASET

Performance Metric	Measured Value
Overall Accuracy	98.5%
Identification Precision	97.9%
Detection Recall	98.2%
Harmonic F1-Score	98.0%
Mean Processing Latency	85 ms
Sustained Frame Rate	11.2 FPS
False Positive Rate	2.1%
False Negative Rate	1.8%

C. Experimental Results

Table I exhibits the capability metrics accomplished by our infrastructure across the comprehensive test dataset.

The outcomes exhibit that the infrastructure accomplishes elevated precision in distinguishing between authorized and unauthorized individuals while sustaining processing velocities compatible with instantaneous functionality. The F1-score of 98.0 percent designates balanced capability across both precision and recall, suggesting the infrastructure effectively minimizes both false affirmatives and false negatives.

The sustained frame rate of 11.2 FPS, while below the camera's native 30 FPS capture rate, proves sufficient for security monitoring applications where human movement occurs over multiple frame intervals. The processing latency of 85 milliseconds enables response times well within human perception thresholds.

D. Comparative Analysis

To contextualize our outcomes, we compare against two alternative methodologies: conventional computer vision techniques (Haar Cascades with Local Binary Pattern Histograms) and an alternative neural learning sequence (MTCNN with FaceNet). Table II exhibits the comparative outcomes across identical test conditions.

TABLE II
COMPARATIVE SYSTEM PERFORMANCE ANALYSIS

Metric	Proposed	Haar+LBPH	MTCNN+FaceNet
Accuracy	98.5%	85.2%	97.9%
Precision	97.9%	88.1%	96.5%
Recall	98.2%	84.5%	97.1%
F1-Score	98.0%	86.3%	96.8%
Latency	85 ms	35 ms	150 ms
Frame Rate	11.2 FPS	28.5 FPS	6.7 FPS

The conventional methodology accomplishes the lowest latency yet exhibits substantially inferior precision, rendering it inappropriate for dependable security implementations where false alarms undermine system utility. The MTCNN with FaceNet sequence delivers robust precision yet at expenditure of significantly augmented latency, culminating in deteriorated instantaneous capability. Our proposed infrastructure accomplishes the highest precision metrics while sustaining latency approximately 43 percent diminished than the MTCNN approach, confirming our architectural construction selections.

The comparative analysis demonstrates that the proposed system occupies an optimal position in the accuracy-speed design space, achieving near-optimal accuracy while maintaining processing speeds significantly better than alternative deep learning approaches.

V. DISCUSSION AND ANALYSIS

A. Architectural Effectiveness and Design Validation

The experimental outcomes confirm our hypothesis that a meticulously constructed cascaded infrastructure can successfully balance competing specifications of processing velocity and analytical precision. The fundamental insight demonstrates that not all frames necessitate expensive analytical operations—by intelligently filtering the processing sequence, computational capabilities can be concentrated where they deliver maximum value.

The multi-threaded execution framework proves indispensable for sustaining operator experience quality. Without concurrent processing, analytical operations would obstruct video capture, culminating in irregular or frozen video transmissions. The executed infrastructure sustains fluid video presentation even during peak analytical burden, with frame delivery consistency exceeding 95 percent across all test scenarios.

The cascaded processing approach effectively implements a computational budget allocation strategy, dedicating expensive operations only to frames warranting detailed analysis. This approach reduces overall computational requirements by approximately 65-70 percent compared to processing all frames through the complete pipeline.

B. Practical Deployment Implications

The exhibited capability characteristics designate the infrastructure as appropriate for implementation in authentic residential environments. The accomplished precision levels substantially surpass specifications for pragmatic utilization, while the processing latency remains within boundaries that facilitate effective instantaneous response capabilities.

Particularly significant proves the infrastructure's functionality on standard consumer equipment without specialized acceleration. This eliminates implementation barriers associated with expensive specialized apparatus, rendering the technology available to characteristic homeowners. The system's ability to operate effectively on CPU-only hardware substantially reduces deployment costs and complexity.

The web-based interface architecture facilitates straightforward integration with existing home network infrastructures without requiring specialized client software installations. Users can access monitoring functionality from any device with a modern web browser, enhancing accessibility and usability.

C. System Limitations and Constraints

Several limitations merit acknowledgment for comprehensive evaluation. First, the infrastructure's capability depends

on sufficient illumination circumstances—severe darkness or extreme backlighting may deteriorate authentication precision. Preliminary testing suggests performance degradation begins at illumination levels below 50 lux, though infrared supplementation could address this limitation.

Second, the current execution processes a singular video transmission; scaling to multiple cameras would necessitate additional optimization of resource allocation and scheduling mechanisms. The single-stream limitation restricts coverage area but simplifies deployment for typical residential applications.

Third, the infrastructure demands initial enrollment of authorized operators through a manual registration procedure requiring capture of multiple facial samples under varying conditions. This enrollment process typically requires 3-5 minutes per user, representing a modest but non-negligible setup overhead.

The processing frame frequency of approximately 11 FPS, while adequate for security observation, descends short of the complete camera capture frequency. This signifies rapid movements may not be captured in every frame, though the cascaded infrastructure guarantees critical incidents are not overlooked due to temporal sampling. The frame skip strategy prioritizes sustained real-time operation over complete temporal coverage.

D. Security and Privacy Considerations

While the infrastructure significantly enhances upon conventional methodologies, it remains not impervious to sophisticated attacks. Presentation attacks utilizing elevated-quality photographs or video replays could potentially deceive the authentication infrastructure. Executing liveness identification would address this vulnerability yet would augment computational specifications. For residential implementations, the current infrastructure delivers sufficient protection for characteristic threat scenarios where sophisticated spoofing attacks remain unlikely.

Privacy considerations warrant attention given the sensitive nature of biometric data. The current implementation stores facial embeddings locally without cloud transmission, addressing fundamental privacy concerns. However, comprehensive deployment should incorporate encryption for stored embeddings and secure communication protocols for remote access scenarios.

The system maintains audit logs of all detection events, supporting forensic analysis while potentially raising data retention privacy questions. Deployment in regulated environments may require implementation of data retention policies and user consent mechanisms.

VI. Conclusions and Future Research Directions

This investigation exhibits that cognitive residential security infrastructures merging rapid object identification with comprehensive biometric authentication can accomplish both elevated precision and instantaneous capability on consumer equipment. The fundamental contribution represents an architectural methodology that strategically distributes computational capabilities through a cascaded processing sequence and concurrent execution framework.

Empirical confirmation validates the infrastructure accomplishes 98.5 percent precision in identity authentication while sustaining processing latencies compatible with instantaneous functionality. Comparative examination exhibits enhanced capability relative to both conventional computer vision techniques and alternative neural learning infrastructures. The demonstrated performance characteristics indicate practical viability for real-world residential security deployments.

The research contributes both theoretical insights into the design of real-time AI systems operating under resource constraints and practical implementation strategies applicable to broader domains beyond residential security. The computational funnel architecture principle generalizes to other applications requiring balance between processing speed and analytical depth.

A. Prospective Investigation Trajectories

Several promising extensions merit exploration in future work:

Multi-Camera Integration: Extending the infrastructure to process multiple video transmissions simultaneously would facilitate comprehensive coverage of larger properties. This necessitates optimization of scheduling and resource distribution mechanisms, potentially leveraging intelligent frame prioritization based on motion detection or region-of-interest analysis.

Behavioral Pattern Analysis: Incorporating temporal examination to identify unusual activity configurations (prolonged loitering, unusual movement trajectories, abnormal temporal patterns) would enhance threat identification capabilities beyond simple presence identification. Machine learning approaches could model normal behavioral patterns and flag deviations.

Edge Computing Implementation: Porting the infrastructure to dedicated edge computing platforms would facilitate completely autonomous functionality without cloud connectivity, addressing privacy concerns associated with distant video transmission. Specialized hardware accelerators could substantially improve processing throughput.

Adaptive Processing Strategies: Executing dynamic adjustment of processing parameters based on current infrastructure burden and threat level could further optimize the effectiveness-precision compromise. Adaptive algorithms could modulate frame sampling rates and detection thresholds based on contextual factors.

Enhanced Authentication Mechanisms: Integrating liveness identification and multi-modal biometrics (merging facial authentication with voice or gait examination) would strengthen protection against sophisticated presentation attacks. Multi-factor approaches could substantially reduce vulnerability to spoofing while maintaining usability.

Federated Learning for Privacy-Preserving Updates: Implementing federated learning mechanisms would enable continuous model improvement across deployed installations while preserving individual privacy through decentralized training approaches.

The exhibited viability of this architectural methodology establishes a foundation for next-generation cognitive security infrastructures that deliver proactive protection rather than passive documentation, contributing to the evolution of intelligent residential environments.

REFERENCES

[1] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Las Vegas, NV, USA, 2016, pp. 779–788.

[2] J. Deng, J. Guo, N. Xue, and S. Zafeiriou, "ArcFace: Additive angular margin loss for deep face recognition," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit.*, Long Beach, CA, USA, 2019, pp. 4690–4699.

[3] K. Zhang, Z. Zhang, Z. Li, and Y. Qiao, "Joint face detection and alignment using multitask cascaded convolutional networks," *IEEE Signal Process. Lett.*, vol. 23, no. 10, pp. 1499–1503, Oct. 2016.

[4] F. Schroff, D. Kalenichenko, and J. Philbin, "FaceNet: A unified embedding for face recognition and clustering," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Boston, MA, USA, 2015, pp. 815–823.

[5] S. H. M. S. M. A. Almalki, S. U. Rehman, and J. M. Aquiar, "An intelligent IoT-based real-time home security system using machine learning," *IEEE Internet Things J.*, vol. 8, no. 12, pp. 9848–9857, June 2021.

[6] Y. Wang, H. Liu, and J. Han, "A lightweight deep learning model for face recognition in smart home environments," *Sensors*, vol. 20, no. 18, p. 5239, Sept. 2020.

[7] InsightFace Project, "InsightFace: 2D and 3D face analysis project," 2023. [Online]. Available: https://insightface.ai/

[8] Ultralytics, "YOLOv8: State-of-the-art object detection," 2023. [Online]. Available: https://github.com/ultralytics/ultralytics

[9] S. Balarabe, "A multi-threading approach for faster OpenCV video streaming in Python," Medium, 2022. [Online]. Available: https://sihabsahariar.medium.com/a-multi-threading-approach-for-faster-opencv-video-streaming-in-python-7324f11dbd2f

[10] H. Wang, Y. Wang, Z. Zhou, X. Ji, D. Gong, J. Zhou, J. Li, and W. Liu, "CosFace: Large margin cosine loss for deep face recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Salt Lake City, UT, USA, 2018, pp. 5265–5274.

[11] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Las Vegas, NV, USA, 2016, pp. 770–778.

[12] J. Guo, J. Deng, Z. Zhou, and S. Zafeiriou, "Sample and computation redistribution for efficient face detection," *arXiv preprint arXiv:2105.04714*, 2021.

[13] G. Bradski, "The OpenCV library," *Dr. Dobb's J. Softw. Tools*, vol. 25, no. 11, pp. 120–123, 2000.

[14] A. Al-Fuqaha, M. Guizani, M. Mohammadi, M. Aledhari, and M. Ayyash, "Internet of Things: A survey on enabling technologies, protocols, and applications," *IEEE Commun. Surveys Tuts.*, vol. 17, no. 4, pp. 2347–2376, Fourthquarter 2015.

[15] M. Grinberg, *Flask Web Development: Developing Web Applications with Python*, 2nd ed. Sebastopol, CA, USA: O'Reilly Media, 2018.

[16] S. R. Kumar, A. K. Ata, et al., "Design and development of an AI based intelligent door for home," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Inform. (ACCAI)*, Bangalore, India, 2023, pp. 1–6.

[17] T. V. Dang, "Smart home management system with face recognition based on ArcFace model in deep convolutional neural network," *J. Robot. Control (JRC)*, vol. 3, no. 6, pp. 754–763, 2022.

[18] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal speed and accuracy of object detection," *arXiv preprint arXiv:2004.10934*, 2020.