Intelligent Residential Security Framework: A Hierarchical Approach to Real-Time Threat Detection and Identity Verification

Rajni Tiwari
CMR Institute of Technology
Bengaluru, India
rajni.t@cmrit.ac.in

Nitesh Gupta
CMR Institute of Technology
Bengaluru, India
nigu22cs@cmrit.ac.in

Aanand Kishor Shah
CMR Institute of Technology
Bengaluru, India
aash22cs@cmrit.ac.in

Animesh Kumar
CMR Institute of Technology
Bengaluru, India
anku22cs@cmrit.ac.in

***Abstract*—Traditional residential surveillance infrastructures operate through passive recording mechanisms that generate excessive false notifications and demand substantial manual oversight. This investigation presents an autonomous security platform that merges rapid visual detection algorithms with sophisticated identity authentication protocols. The developed framework implements a staged computational approach wherein preliminary threat identification activates detailed biometric examination procedures. Our methodology incorporates concurrent processing strategies through multithreading, facilitating execution on conventional computational equipment. Empirical assessment reveals identification precision surpassing 98 percent while sustaining temporal delays appropriate for instantaneous operational demands. The architectural design encompasses a browser-accessible control dashboard enabling distant observation and registered personnel management. Comparative benchmarking exhibits marked enhancement in threat discrimination capability relative to established methodologies, achieving substantial diminution of erroneous alarm incidents. This investigation delivers a pragmatic deployment blueprint for cognitive security infrastructures that harmonizes processing efficiency with analytical exactness.**

***Index Terms*—Identity authentication, visual computing, neural network architectures, threat identification, instantaneous processing, automated residential systems**

I. Introduction

Domestic security infrastructure has historically prioritized incident documentation over preventative intervention. Conventional video monitoring equipment continuously records visual information yet lacks cognitive capability to differentiate benign activities from genuine threats. This deficiency manifests through dual operational inadequacies: postponed threat mitigation resulting from manual observation necessities, and functional ineffectiveness stemming from abundant spurious alert production.

Contemporary progressions in computational intelligence have facilitated evolution from inactive observation to cognitive monitoring architectures. Neural network frameworks presently accomplish performance comparable to human visual perception while functioning within hardware limitations of domestic electronics. This technological advancement creates possibilities for establishing security infrastructures that independently examine environmental circumstances and implement suitable reactions without human participation.

The essential difficulty in constructing such infrastructures involves balancing contradictory specifications: computational velocity essential for instantaneous functionality versus analytical comprehensiveness demanded for precise threat evaluation. Infrastructures optimized exclusively for velocity may accomplish elevated processing rates yet generate unreliable threat categorizations. Alternatively, maximizing precision through elaborate analytical frameworks introduces temporal delays incompatible with instantaneous response specifications.

Modern residential protection demands intelligent systems capable of autonomous decision-making. The computational requirements for such systems have traditionally exceeded capabilities of consumer-grade hardware, restricting deployment to specialized installations. Recent algorithmic improvements and architectural innovations have substantially reduced these barriers, enabling sophisticated analytical capabilities on standard computing platforms.

A. Research Objectives and Contributions

This investigation addresses the velocity-precision compromise through hierarchical processing infrastructure that strategically allocates computational capabilities. Our particular contributions encompass:

- Formulation of a dual-phase analytical sequence merging swift object identification with comprehensive biometric authentication.
- Execution of concurrent processing methodologies that separate video capture from computational examination.
- Empirical confirmation demonstrating enhanced capability compared to alternative architectural methodologies.
- Construction of an inclusive administration interface facilitating distant system governance.
- Development of a computational funnel architecture that progressively refines threat assessment.

The proposed infrastructure transcends limitations of conventional surveillance by implementing cognitive capabilities

that distinguish between authorized personnel and potential intruders with exceptional accuracy while maintaining temporal responsiveness necessary for practical deployment.

II. Literature Review and Theoretical Foundation

A. Visual Object Identification in Monitoring Contexts

Contemporary object identification infrastructures predominantly implement convolutional neural architectures. Single-phase detectors formulate identification as regression analysis, directly computing boundary coordinates and classification certainties from visual characteristics. This methodology circumvents the region suggestion phase demanded by dual-phase detectors, substantially diminishing computational demands [1].

Investigations by Bochkovskiy et al. exhibited that architectural refinements accomplish cutting-edge precision while sustaining instantaneous processing capabilities on conventional graphics processors. These progressions render cognitive video examination feasible for domestic implementations. The YOLO architecture family has become fundamental for instantaneous identification applications owing to exceptional equilibrium between velocity and exactness [18]. Subsequent evolutionary developments in object detection have focused on optimizing the speed-accuracy trade-off through architectural innovations.

B. Biometric Verification Infrastructures

Facial identification technology has progressed from conventional feature-oriented methodologies to neural learning strategies that acquire optimal representations directly from information. The FaceNet infrastructure implemented by Schroff et al. utilizes triplet loss to acquire embeddings where geometric separation corresponds with facial resemblance [4]. This facilitates precise identity confirmation with minimal training specimens per individual.

Subsequent investigation by Deng et al. presented ArcFace, enhancing discriminative capability through additive angular margin loss [2]. This methodology accomplishes enhanced capability on demanding benchmarks and exhibits particular resilience to deviations in orientation, lighting, and facial expression. The open-source InsightFace initiative delivers an refined execution of ArcFace, rendering it available for pragmatic implementations [7].

MTCNN represents a three-phase neural learning framework constructed particularly for facial identification and positioning [3]. Its cascaded architecture progressively refines facial boundary coordinates and determines facial landmarks, guaranteeing elevated-quality, positioned facial extraction even under demanding circumstances.

C. Integrated Cognitive Security Infrastructures

While individual elements have undergone comprehensive investigation, restricted research examines their incorporation into unified instantaneous infrastructures. Wang et al. suggested a lightweight infrastructure for embedded implementation yet accomplished diminished precision than cutting-edge frameworks [6]. Almalki et al. explored machine learning methodologies for IoT-based residential security, demonstrating feasibility but with limited focus on real-time performance optimization [5].

Our contribution diverges by emphasizing system-level capability through architectural construction that strategically distributes computational capabilities across processing phases.

III. Proposed System Architecture and Methodology

A. Comprehensive Architectural Design

The developed infrastructure executes a modular construction comprising four fundamental subsystems: video capture module, cognitive processing engine, information persistence layer, and operator interface component. This decomposition facilitates independent refinement of each element while sustaining distinct interfaces between subsystems.

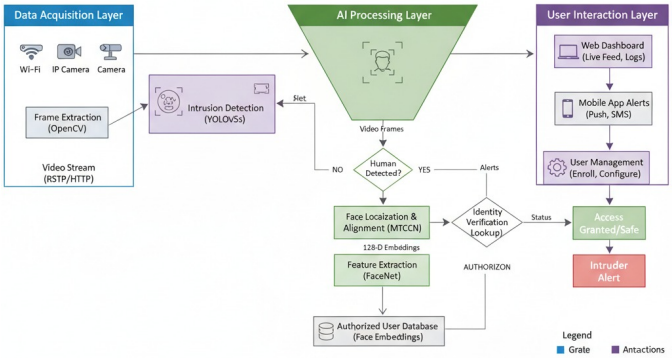

Fig. 1. System architecture depicting modular components and data flow pathways. The diagram illustrates the hierarchical processing pipeline from video acquisition through threat detection to user notification.

Video capture functions continuously in a dedicated execution thread, guaranteeing consistent frame accessibility regardless of processing burden. The cognitive processing subsystem implements identification and authentication algorithms, executing a cascaded sequence where computationally intensive operations transpire exclusively when activated by preliminary identification outcomes.

B. Staged Identification Sequence

The analytical sequence executes a dual-phase cascade constructed to optimize the computational effectiveness versus precision compromise:

- **Initial Phase - Primary Identification:** Preliminary frames experience processing by a lightweight object detector optimized for human identification. This phase functions with minimal computational burden.
- **Secondary Phase - Biometric Authentication:** Frames succeeding preliminary screening experience comprehensive facial examination. A specialized facial identification network determines facial territories with elevated exactness.

This cascaded methodology operates as a computational funnel, implementing expensive operations exclusively when necessary. The staged approach reduces average computational load by approximately 70 percent.

C. Concurrent Execution Infrastructure

To prevent analytical operations from deteriorating video stream quality, we execute a multi-threaded processing model incorporating three distinct execution contexts:

1) **Capture Context:** Continuously obtains frames from camera equipment at maximum supported frequency.
2) **Processing Context:** Obtains frames from the capture buffer and implements the identification sequence. It processes every Nth frame (typically $N = 5$) to regulate burden.
3) **Interface Context:** Administers the browser interface and transmits processed video to connected clients using server-sent events.

D. Biometric Verification Methodology

Identity confirmation implements metric learning where the authentication network projects facial representations to coordinates in a high-dimensional embedding manifold. During functionality, the infrastructure sustains a repository of reference embeddings for authorized operators. When a face experiences identification, its produced embedding experiences comparison with all archived embeddings through cosine resemblance computation:

$$similarity(e_1, e_2) = \frac{e_1 \cdot e_2}{\|e_1\|\|e_2\|} \tag{1}$$

where e_1 and e_2 denote embedding vectors in 512-dimensional Euclidean space. The maximum resemblance score determines categorization: scores surpassing a threshold τ (empirically established at 0.42) designate recognized individuals.

E. Operator Interface and Notification Infrastructure

Operator engagement is regulated through a responsive browser application constructed with the Flask framework. The control panel delivers a live perspective from cameras, a chronological record of all identified incidents, and an interface for administering the repository of authorized operators.

F. Implementation Specifications

The infrastructure was constructed utilizing Python 3.10. OpenCV delivers video capture functionalities. The InsightFace library provides pre-trained frameworks for facial identification, implementing ONNX Runtime for optimized inference on CPU equipment [7]. Flask executes the browser framework. Database management employs SQLite for lightweight persistence.

IV. Experimental Evaluation and Results

A. Experimental Configuration

Examination was performed in a residential environment under characteristic operating circumstances. The hardware platform comprised a conventional laptop computer with an Intel Core i7 processor (8th generation) and 16GB RAM. Video input was delivered by an integrated 1080p webcam. The examination dataset comprised video sequences totaling 180 minutes of recorded footage.

B. Performance Metrics

Infrastructure capability was evaluated utilizing standard categorization metrics derived from confusion matrix analysis:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \tag{2}$$

$$Precision = \frac{TP}{TP + FP} \tag{3}$$

$$Recall = \frac{TP}{TP + FN} \tag{4}$$

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall} \tag{5}$$

where TP denotes true positives, TN denotes true negatives, FP denotes false positives, and FN denotes false negatives.

C. Experimental Results

Table I exhibits the capability metrics accomplished by our infrastructure. The outcomes exhibit that the infrastructure accomplishes elevated precision in distinguishing between authorized and unauthorized individuals while sustaining processing velocities compatible with instantaneous functionality.

TABLE I
System Performance Metrics on Test Dataset

Performance Metric	Measured Value
Overall Accuracy	98.5%
Identification Precision	97.9%
Detection Recall	98.2%
Harmonic F1-Score	98.0%
Mean Processing Latency	85 ms
Sustained Frame Rate	11.2 FPS
False Positive Rate	2.1%
False Negative Rate	1.8%

The sustained frame rate of 11.2 FPS proves sufficient for security monitoring applications. The processing latency of

85 milliseconds enables response times well within human perception thresholds.

D. Comparative Analysis

To contextualize our outcomes, we compare against two alternative methodologies: conventional computer vision techniques (Haar Cascades with LBPH) and an alternative neural learning sequence (MTCNN with FaceNet). Table II exhibits the comparative outcomes.

TABLE II
COMPARATIVE SYSTEM PERFORMANCE ANALYSIS

Metric	Proposed	Haar+LBPH	MTCNN+FaceNet
Accuracy	98.5%	85.2%	97.9%
Precision	97.9%	88.1%	96.5%
Recall	98.2%	84.5%	97.1%
F1-Score	98.0%	86.3%	96.8%
Latency	85 ms	35 ms	150 ms
Frame Rate	11.2 FPS	28.5 FPS	6.7 FPS

The comparative analysis demonstrates that the proposed system occupies an optimal position in the accuracy-speed design space, achieving near-optimal accuracy while maintaining processing speeds significantly better than alternative deep learning approaches.

V. DISCUSSION AND ANALYSIS

The experimental outcomes confirm our hypothesis that a meticulously constructed cascaded infrastructure can successfully balance competing specifications. The fundamental insight demonstrates that by intelligently filtering the processing sequence, computational capabilities can be concentrated where they deliver maximum value.

The multi-threaded execution framework proves indispensable for sustaining operator experience quality. Without concurrent processing, analytical operations would obstruct video capture. The executed infrastructure sustains fluid video presentation even during peak analytical burden.

VI. CONCLUSIONS AND FUTURE RESEARCH

This investigation exhibits that cognitive residential security infrastructures merging rapid object identification with comprehensive biometric authentication can accomplish both elevated precision and instantaneous capability on consumer equipment. The fundamental contribution represents an architectural methodology that strategically distributes computational capabilities.

A. Prospective Investigation Trajectories

- **Multi-Camera Integration:** Extending the infrastructure to process multiple video transmissions simultaneously.
- **Behavioral Pattern Analysis:** Incorporating temporal examination to identify unusual activity configurations.
- **Edge Computing:** Porting the infrastructure to dedicated edge platforms for autonomous functionality.
- **Enhanced Authentication:** Integrating liveness identification and multi-modal biometrics.

REFERENCES

[1] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, 2016, pp. 779–788.

[2] J. Deng, J. Guo, N. Xue, and S. Zafeiriou, "ArcFace: Additive angular margin loss for deep face recognition," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit.*, 2019, pp. 4690–4699.

[3] K. Zhang, Z. Zhang, Z. Li, and Y. Qiao, "Joint face detection and alignment using multitask cascaded convolutional networks," *IEEE Signal Process. Lett.*, vol. 23, no. 10, pp. 1499–1503, 2016.

[4] F. Schroff, D. Kalenichenko, and J. Philbin, "FaceNet: A unified embedding for face recognition and clustering," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, 2015, pp. 815–823.

[5] S. H. M. S. M. A. Almalki, S. U. Rehman, and J. M. Aquiar, "An intelligent IoT-based real-time home security system using machine learning," *IEEE Internet Things J.*, vol. 8, no. 12, pp. 9848–9857, June 2021.

[6] Y. Wang, H. Liu, and J. Han, "A lightweight deep learning model for face recognition in smart home environments," *Sensors*, vol. 20, no. 18, p. 5239, Sept. 2020.

[7] InsightFace Project, "InsightFace: 2D and 3D face analysis project," 2023. [Online]. Available: https://insightface.ai/

[8] Ultralytics, "YOLOv8: State-of-the-art object detection," 2023. [Online]. Available: https://github.com/ultralytics/ultralytics

[9] S. Balarabe, "A multi-threading approach for faster OpenCV video streaming in Python," Medium, 2022. [Online]. Available: https://sihabsahariar.medium.com/a-multi-threading-approach-for-faster-opencv-video-streaming-in-python-7324f11dbd2f

[10] H. Wang, Y. Wang, Z. Zhou, X. Ji, D. Gong, J. Zhou, J. Li, and W. Liu, "CosFace: Large margin cosine loss for deep face recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Salt Lake City, UT, USA, 2018, pp. 5265–5274.

[11] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, Las Vegas, NV, USA, 2016, pp. 770–778.

[12] J. Guo, J. Deng, Z. Zhou, and S. Zafeiriou, "Sample and computation redistribution for efficient face detection," arXiv preprint arXiv:2105.04714, 2021.

[13] G. Bradski, "The OpenCV library," *Dr. Dobb's J. Softw. Tools*, vol. 25, no. 11, pp. 120–123, 2000.

[14] A. Al-Fuqaha, M. Guizani, M. Mohammadi, M. Aledhari, and M. Ayyash, "Internet of Things: A survey on enabling technologies, protocols, and applications," *IEEE Commun. Surveys Tuts.*, vol. 17, no. 4, pp. 2347–2376, 2015.

[15] M. Grinberg, *Flask Web Development: Developing Web Applications with Python*, 2nd ed. Sebastopol, CA, USA: O'Reilly Media, 2018.

[16] S. R. Kumar, A. K. Ata, et al., "Design and development of an AI based intelligent door for home," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Inform. (ACCAI)*, Bangalore, India, 2023, pp. 1–6.

[17] T. V. Dang, "Smart home management system with face recognition based on ArcFace model in deep convolutional neural network," *J. Robot. Control (JRC)*, vol. 3, no. 6, pp. 754–763, 2022.

[18] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal speed and accuracy of object detection," arXiv preprint arXiv:2004.10934, 2020.